

# Apache Sqoop

BAS Academy


# Agenda

- ▶ About Scoop
- ▶ Sqoop Import
- ▶ Incremental Import
- ▶ Sqoop Commands and Arguments
- ▶ Sqoop Export
- ▶ Integration with Ecosystem
- ▶ Hands On

# About Sqoop

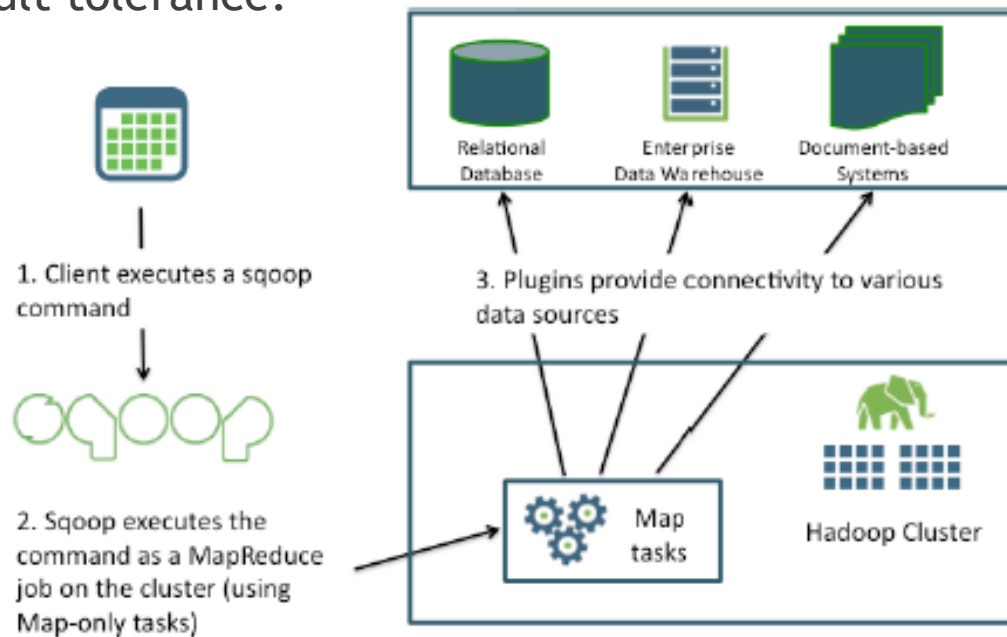
# What is Sqoop

- ▶ **Sqoop:** “SQL to Hadoop and Hadoop to SQL”
- ▶ Sqoop is a tool designed to transfer data between Hadoop and relational database servers.
- ▶ It is used to import data from relational databases such as MySQL, Oracle to Hadoop HDFS, and export from Hadoop file system to relational databases.
- ▶ It is provided by the Apache Software Foundation.



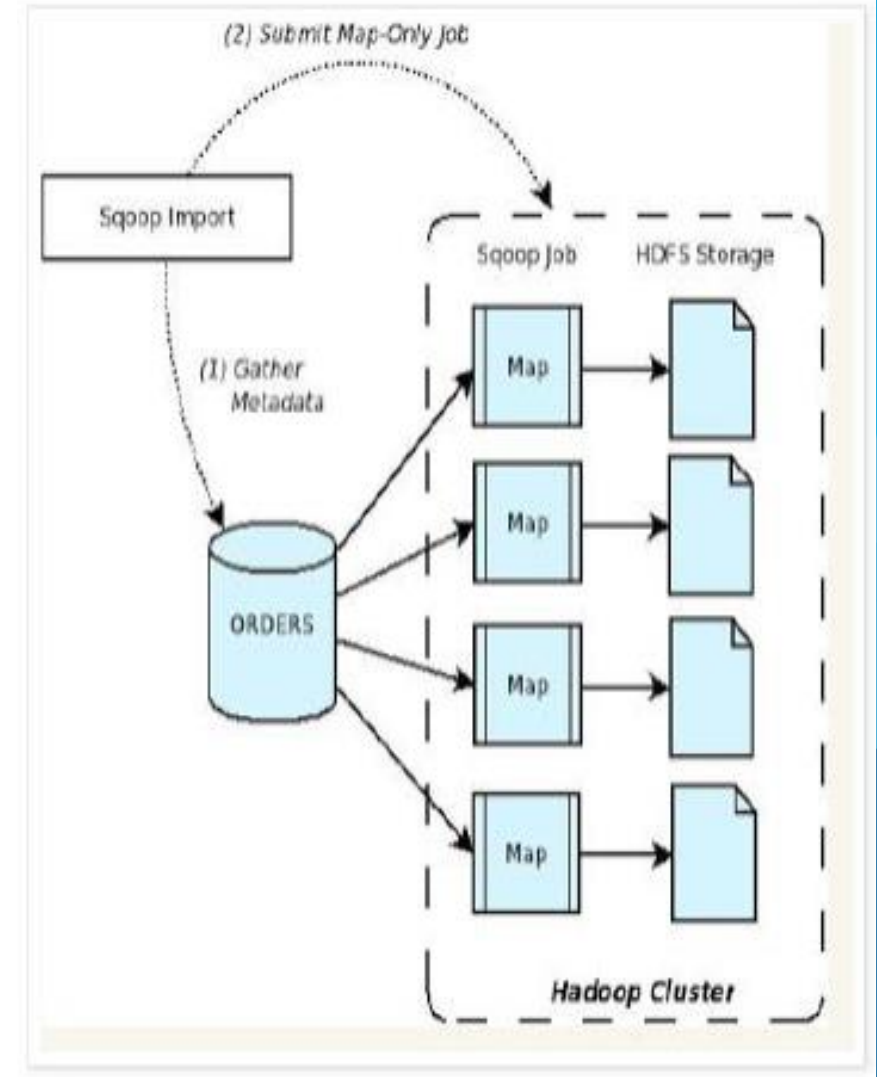
# Sqoop Architecture

- ▶ Sqoop uses a connector-based architecture that supports plugins that provide connectivity to additional external systems.
- ▶ Sqoop uses MapReduce to distribute its work across the Hadoop cluster
- ▶ Using MapReduce to perform Sqoop commands provides parallel operation as well as fault tolerance.



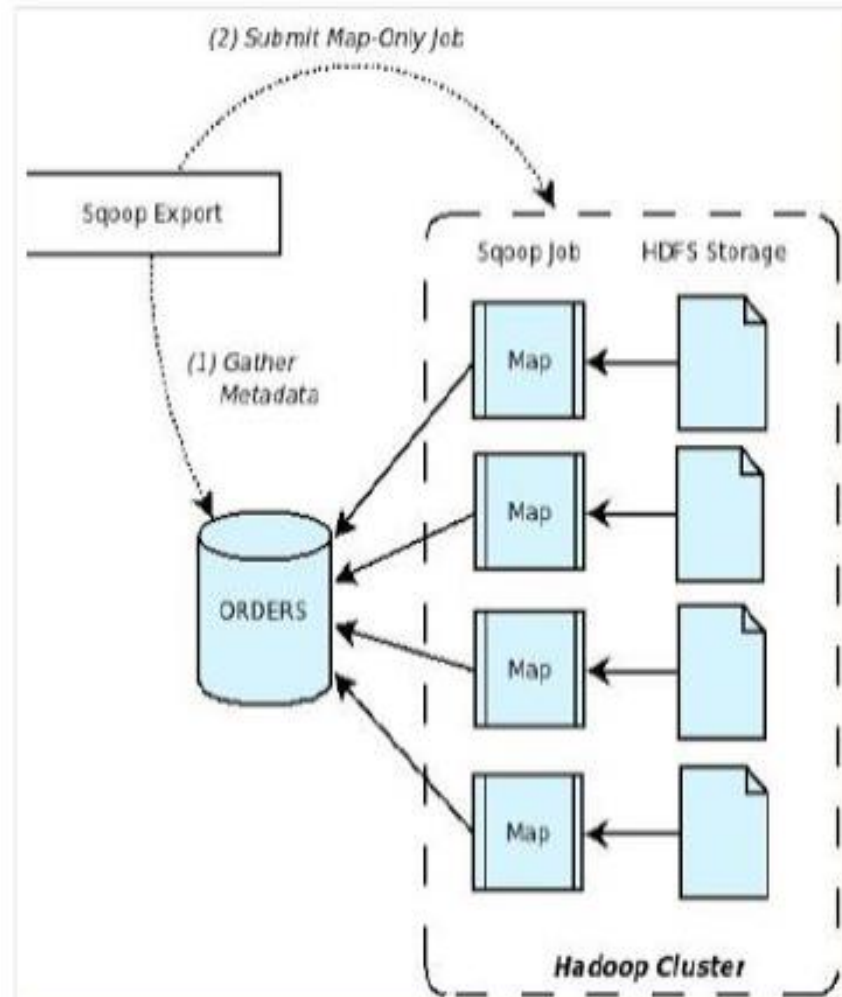
# Sqoop Import

- ▶ The input to the import process is a database table
- ▶ Sqoop will read the table row-by-row into HDFS.
- ▶ The output of this import process is a set of files containing a copy of the imported table.
- ▶ The import process is performed in parallel. For this reason, the output will be in multiple files.
- ▶ These files may be delimited text files (for example, with commas or tabs separating each field) or binary Avro or SequenceFiles



# Sqoop Export

- ▶ After manipulating the imported records (for example, with MapReduce or Hive) you may have a result data set which you can then export back to the relational database.
- ▶ The export tool exports a set of files from HDFS back to an RDBMS.
- ▶ The files given as input to Sqoop contain records, which are called as rows in table.
- ▶ Those are read in parallel and parsed into a set of records and inserted them as new rows in a target database table



# Sqoop Import



# Importing Tables

The import command has the following requirements:

- ▶ Must specify a connect string using the `--connect` argument
- ▶ Can include credentials in the connect string, using the `--username` and `--password` arguments
- ▶ Must specify either a table to import using `--table` or the result of an SQL query using `-query`
- ▶ The default number of map tasks for Sqoop is four, so the result of default import will be in four files

```
sqoop import
--connect jdbc:mysql://host/nyse
--table StockPrices
--target-dir /data/stockprice/
--as-textfile
```

# File Format

Sqoop supports Text and Binary files.

- ▶ Text

Import into text file using the `--as-textfile` parameter:

- ▶ Binary formats are a natural fit for storing binary values like images or PDF documents
- ▶ To access the binary data, you need to implement extra functionality or load special libraries in your application

## Binary Files:

- ▶ Avro

Import into Avro file by specifying the `--as-avrodatafile` parameter

- ▶ SequenceFile

Import into SequenceFile using the `--as-sequencefile` parameter:

# Import All Tables

- ▶ Sqoop has the feature to import all the tables at once
- ▶ Tables will be imported in sequential order

```
sqoop import-all-tables \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop
```

- ▶ Option to exclude few tables

```
sqoop import-all-tables \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop \  
  --exclude-tables cities,countries
```

# Freeform SQL

- ▶ Instead of using table import, use free-form query import.
- ▶ Use the --query argument to specify which rows to select from a table.
- ▶ Sqoop will not use the database catalog to fetch the metadata

```
sqoop import
--connect jdbc:mysql://host/nyse
--query "SELECT * FROM StockPrices s
WHERE s.Volume >= 1000000
AND \$CONDITIONS"
--target-dir /data/highvolume/
--as-textfile
--split-by StockSymbol
```

- ▶ --split-by parameter with the column is used for slicing the data into multiple parallel tasks.
- ▶ Using --query is limited to simple queries

# Incremental Import

# Incremental Import

- ▶ Sqoop provides an incremental import mode which can be used to retrieve only rows newer than some previously-imported set of rows.

Sqoop supports two types of incremental imports:

- ▶ append

For importing the newly created rows to the existing data

- ▶ lastmodified

Used when existing rows needs to be updated

- ▶ --incremental argument is used for incremental import

# Incremental Import - Append

- ▶ Incremental import in append mode will allow you to transfer only the newly created rows.

Incremental import also requires two additional parameters:

- ▶ `--check-column` indicates a column name that should be checked for newly appended data
- ▶ `-last-value` contains the last value that successfully imported into Hadoop.
- ▶ Sqoop when running in incremental mode, always prints out the value of the last imported row.

```
sqoop import \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop \  
  --table visits \  
  --incremental append \  
  --check-column id \  
  --last-value 1
```

# Incremental Import - Lastmodified

- ▶ This is used for mutable data. The data which is getting changed
- ▶ Use the lastmodified mode instead of the append mode.
- ▶ The incremental mode lastmodified requires a column holding a date value (suitable types are date, time, datetime, and timestamp) containing information as to when each row was last updated.

```
sqoop import \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop \  
  --table visits \  
  --incremental lastmodified \  
  --check-column last_update_date \  
  --last-value "2013-05-22 01:01:01"
```



# Sqoop Jobs

- ▶ The Sqoop metastore allows you to retain your job definitions and to easily run them anytime.
- ▶ Each saved job has a logical name that is used for referencing.
- ▶ To create a sqoop job:

```
sqoop job \  
  --create visits \  
  -- \  
  import \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop \  
  --table visits \  
  --incremental append \  
  --check-column id \  
  --last-value 0
```

- ▶ List all retained jobs using the --list parameter

```
sqoop job -list
```

- ▶ View content of the saved job definitions using the -show parameter

```
sqoop job --show visits
```

- ▶ Remove the old job definitions that are no longer needed with the -delete parameter

```
sqoop job --delete visits
```

# Sqoop Commands and Arguments

# Sqoop Commands

Tool specific arguments start with two dashes (--)

```
Available commands:
codegen          Generate code to interact with database records
create-hive-table Import a table definition into Hive
eval            Evaluate a SQL statement and display the results
export          Export an HDFS directory to a database table
help            List available commands
import          Import a table from a database to HDFS
import-all-tables Import tables from a database to HDFS
list-databases  List available databases on a server
list-tables     List available tables in a database
version         Display version information
```

See 'sqoop help COMMAND' for information on a specific command.

Hadoop specific arguments are preceded by single dash character (-)

```
Common arguments:
--connect <jdbc-uri>      Specify JDBC connect string
--connect-manager <jdbc-uri> Specify connection manager class to use
--driver <class-name>    Manually specify JDBC driver class to use
--hadoop-home <dir>      Override $HADOOP_HOME
--help                   Print usage instructions
-P                       Read password from console
--password <password>    Set authentication password
--username <username>    Set authentication username
--verbose                Print more information while working
```

# Sqoop Import Control Args

Argument	Description
<code>--append</code>	Append data to an existing dataset in HDFS
<code>--as-avrodatafile</code>	Imports data to Avro Data Files
<code>--as-sequencefile</code>	Imports data to SequenceFiles
<code>--as-textfile</code>	Imports data as plain text (default)
<code>--boundary-query &lt;statement&gt;</code>	Boundary query to use for creating splits
<code>--columns &lt;col,col,col...&gt;</code>	Columns to import from table
<code>--direct</code>	Use direct import fast path
<code>--direct-split-size &lt;n&gt;</code>	Split the input stream every <i>n</i> bytes when importing in direct mode
<code>--inline-lob-limit &lt;n&gt;</code>	Set the maximum size for an inline LOB
<code>-m, --num-mappers &lt;n&gt;</code>	Use <i>n</i> map tasks to import in parallel
<code>-e, --query &lt;statement&gt;</code>	Import the results of <i>statement</i> .
<code>--split-by &lt;column-name&gt;</code>	Column of the table used to split work units
<code>--table &lt;table-name&gt;</code>	Table to read
<code>--target-dir &lt;dir&gt;</code>	HDFS destination dir
<code>--warehouse-dir &lt;dir&gt;</code>	HDFS parent for table destination
<code>--where &lt;where clause&gt;</code>	WHERE clause to use during import
<code>-z, --compress</code>	Enable compression
<code>--compression-codec &lt;c&gt;</code>	Use Hadoop codec (default gzip)
<code>--null-string &lt;null-string&gt;</code>	The string to be written for a null value for string columns
<code>--null-non-string &lt;null-string&gt;</code>	The string to be written for a null value for non-string columns

The `--null-string` and `--null-non-string` arguments are optional. If not specified, then the string "null" will be used.

# Sqoop Export

# Sqoop Export

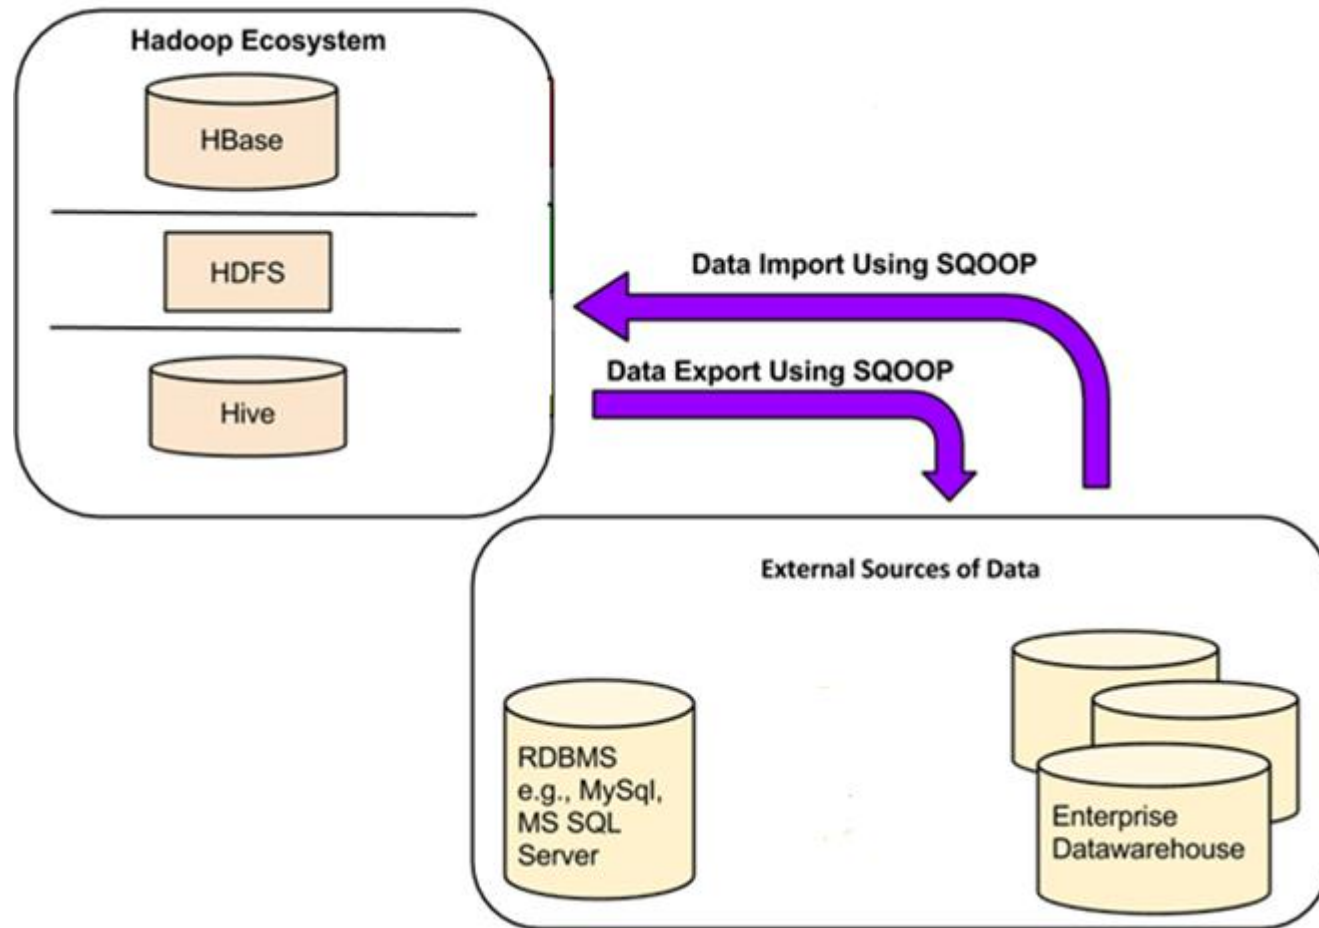
- ▶ Sqoop's export process will read a set of delimited text files from HDFS in parallel, parse them into records, and insert them as new rows in a target database table.
- ▶ Data can be exported from Hadoop to the database on an iterative basis.
- ▶ The only requirement is that there not be any constraint violations when performing the INSERT statements
- ▶ With insert mode, records exported by Sqoop are appended to the end of the target table
- ▶ With Update mode, it works as if exists update else insert
- ▶ With Call mode, it makes stored procedure call

```
sqoop export
--connect jdbc:mysql://host/mylogs
--table LogData
--export-dir /data/logfiles/
--input-fields-terminated-by "\t"
```

# Integration with Ecosystem

# Integration with Ecosystem

- Sqoop can be integrated with the rest of the Hadoop Ecosystem





# Sqoop with Hive

- ▶ Sqoop can import your data directly into Hive.
- ▶ Add the parameter `--hive-import` to your command to enable it
- ▶ Example

```
$ sqoop create-hive-table
```

```
--connect jdbc:mysql://localhost:3306/flights  
--table Flights  
--username root  
--password cloudera
```

```
sqoop import -m 1 --connect jdbc:mysql://localhost:3306/flights --table FLIGHTS --username root --password cloudera  
--hive-import
```

# Sqoop with HBase

To enable import into HBase, there are two additional parameters:

- ▶ **--hbase-table**

Specifies the name of the table in HBase to which you want to import your data.

- ▶ **--column-family**

Specifies into which column family Sqoop will import your table's data.

```
sqoop import \  
  --connect jdbc:mysql://mysql.example.com/sqoop \  
  --username sqoop \  
  --password sqoop \  
  --table cities \  
  --hbase-table cities \  
  --column-family world
```

# Sqoop with Oozie

- ▶ Sqoop jobs can be scheduled in Oozie
- ▶ Oozie includes special Sqoop actions that you can use to call Sqoop in your workflow.

```
<workflow-app name="sqoop-workflow" xmlns="uri:oozie:workflow:0.1">  
  ...  
  <action name="sqoop-action">  
    <sqoop xmlns="uri:oozie:sqoop-action:0.2">
```

- ▶ <command> to list all the parameters,

```
<command>import --table cities --username sqoop --password sqoop ...</command>
```



Hands On

# Sqoop Hands On

1. **Import all columns, filter rows using where clause**
2. **Use Split by option the query**
3. **Write a sqoop command to perform Incremental imports**
4. **Use eval function to display the sqoop import result on the console**



# Thank You

Keerthiga Barathan